
Pokemon Battling Agents: Implementation and Discussion of Various RL algorithms

Arzaan Singh
arzaan@tulane.edu

Ian Kreger
ikreger@tulane.edu

Michael Weild
mweild@tulane.edu

1

Source code repository

Abstract

2 Competitive Pokémon battles pose a challenging reinforcement learning environment due to partial observability, stochastic transitions, and large action spaces.
3 We implement and compare several RL agents, Tabular Q-Learning, Hierarchical
4 Q-Learning, Linear SARSA, and Deep Q-Networks, using the poke-env interface
5 for Pokémon Showdown. Our experiments show that Tabular Q-Learning signifi-
6 cantly outperforms all other implementations, achieving a 98% win rate against
7 the random baseline and over 56% versus the rule-based MaxBasePower agent.
8 Hierarchical and deep agents exhibit strategic promise but suffer from instability
9 and reward exploitation. These results suggest that carefully engineered tabular
10 methods remain highly competitive in constrained Gen 1 battle settings.

12 **1 Introduction**

13 Competitive Pokémon battles pose a sequential decision-making challenge involving uncertainty,
14 opponent modeling, and long-term planning. Each turn, an agent selects an action, either attacking
15 or switching, based on a partially observed game state that includes type matchups, move effects,
16 and potential counterplay. This structure naturally defines the problem as a Markov Decision
17 Process (MDP), making Pokémon an appealing testbed for reinforcement learning (RL) algorithms
18 in adversarial settings.

19 Pokémon Showdown, along with the poke-env Python interface, enables programmatic access
20 to the battle environment and facilitates scalable training of autonomous agents. We design and
21 compare multiple RL-based Pokémon battlers to understand how different learning paradigms
22 influence performance and convergence. We implement Tabular Q-Learning, Linear SARSA, Deep
23 Q-Networks (DQN), and Hierarchical Q-Learning, demonstrating that simpler tabular approaches
24 can outperform more complex methods in the constrained dynamics of Generation 1 gameplay.

25 **2 Related work**

26 Competitive Pokémon battling has recently transitioned from a community hobby to a structured
27 research domain for evaluating reinforcement learning agents. The PokéAGENT competition for-
28 malized this setting by providing standardized battle environments, leaderboards, and reproducible
29 evaluation protocols [Karten et al., 2025]. Early academic efforts, such as *Optimal Battle Strategy in*
30 *Pokémon using Reinforcement Learning*, demonstrated that Generation 1 battles could be modeled as
31 a Markov Decision Process and explored Q-learning for move selection and tactical planning [Kalose
32 et al., 2018]. Metamon extended this paradigm by introducing scalable adversarial self-play, enabling
33 agents to train over thousands of encounters and reveal emergent counter-strategies [Grigsby, 2024],
34 although these systems remained far from consistently surpassing strong human players.

35 More sophisticated pipelines have since been proposed. An MIT thesis reframed battle optimization
 36 using a hybrid Monte Carlo Tree Search and RL architecture to support long-horizon reasoning
 37 and opponent modeling [Wang, 2024]. These works collectively highlight Pokémon battles as a
 38 challenging RL benchmark due to high branching factors, delayed rewards, and partial observability.
 39 Frameworks such as `poke-env` have facilitated reproducible experimentation by offering state
 40 extraction, action encoding, and Pokémon Showdown connectivity. However, these libraries serve
 41 primarily as infrastructure rather than sources of advanced agents. Recent efforts like PokeChamp
 42 explore Bayesian priors, large language models, and hybrid optimization to improve generalization
 43 across opponents and formats [Kannan et al., 2025], though performance remains unstable and lacks
 44 consensus on best practices.

45 3 Problem formulation

46 We model competitive Pokémon battles as a finite-horizon Markov Decision Process (MDP) defined
 47 by the tuple $\langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle$, where \mathcal{S} is the state space, \mathcal{A} the set of legal actions, \mathcal{P} the transition
 48 dynamics induced by the Pokémon Showdown simulator, \mathcal{R} the reward function, and γ the discount
 49 factor. A battle terminates when all Pokémon on one team faint, yielding a binary outcome, or when
 50 a player exceeds the 10-second turn timer. The environment is adversarial, partially observable, and
 51 combinatorial, requiring the agent to reason about type matchups, switching, move accuracy, status
 52 effects, and long-term board control under stochastic Generation 1 mechanics.

53 3.1 State space

54 Each state s_t is a feature vector extracted from `poke-env`. Depending on the implementation, features
 55 include the identity of active Pokémon, remaining HP (bucketed), stat boosts, type relationships,
 56 status conditions (e.g., SLP, FRZ), and move metadata such as accuracy and base power.
 57 In tabular Q-learning, we represent states using discrete encodings of key battle attributes (e.g.,
 58 Pokémon ID, move ID). This yields a 6-dimensional state vector:

$$s_t = \langle \text{my_hp}, \text{opp_hp}, \text{my_type}, \text{opp_type}, \text{opp_status}, \text{im_faster} \rangle.$$

59 Our hierarchical Q-learning agent decomposes the decision process into a master agent that selects
 60 among five macro-actions (four moves and a switch) and a subagent that selects which Pokémon to
 61 switch into. Their state representations are:

$$s_t^{\text{master}} = \langle \text{self_mon}, \text{opp_mon}, \text{self_hp_bucket}, \text{opp_hp_bucket} \rangle,$$

$$s_t^{\text{sub}} = \langle \text{party_pokemon}, \text{opp_pokemon} \rangle.$$

63 For function approximation methods (DQN and Linear SARSA), we replaced discrete encodings
 64 with continuous features such as normalized HP, one-hot move types, damage estimates, and type-
 65 effectiveness scores. Feature dimensionality ranged from 21 to over 600, with hidden layers of
 66 128–512 units in DQN variants. While Generation 1 constraints keep these representations tractable,
 67 the resulting state spaces are still large enough to prevent full convergence.

68 3.2 Action space

69 At each turn, the agent selects $a_t \in \mathcal{A}(s_t)$, corresponding to one of the available moves (up to four)
 70 or a switch into any unfainted party member. Because available actions depend on team state and
 71 game progression, $\mathcal{A}(s_t)$ is non-stationary, a property that makes the domain more challenging than
 72 fixed-action benchmarks such as Atari. This also explains why hand-engineered bots like `maxbp`
 73 perform competitively in Generation 1 but fail to generalize beyond their encoded heuristics.

74 3.3 Reward structure

75 We initially employed a sparse terminal reward:

$$R(s_t, a_t) = \begin{cases} +1, & \text{win}, \\ -1, & \text{loss}. \end{cases}$$

76 However, the absence of intermediate feedback resulted in slow learning and brittle policy discovery.
 77 To accelerate credit assignment, we introduced reward shaping:

$$R(s_t, a_t) = \begin{cases} +100, & \text{win,} \\ -100, & \text{loss,} \\ +20, & \text{opponent fainted,} \\ -20, & \text{self fainted,} \\ 0.5 \cdot \Delta\text{HP}, & \text{damage dealt,} \\ -0.5 \cdot \Delta\text{HP}, & \text{damage taken.} \end{cases}$$

78 While effective, this formulation introduced reward-hacking behaviors such as repeated healing or
 79 status moves that accrued points without improving board position. To mitigate this, we added
 80 auxiliary penalties and bonuses:

$$R'(s_t, a_t) = \begin{cases} -30, & \text{repeat move more than three times,} \\ +10, & \text{inflict a new status.} \end{cases}$$

81 3.4 Objective

82 The agent seeks a policy $\pi(a | s)$ maximizing the expected discounted return:

$$J(\pi) = \mathbb{E}_\pi \left[\sum_{t=0}^T \gamma^t r_t \right],$$

83 where T is the battle horizon. All Q-learning and SARSA variants were trained using online temporal-
 84 difference bootstrapping, while DQN replaces tabular updates with a neural value approximator
 85 trained from replayed transitions. Since \mathcal{P} is unknown and opponent behavior is stochastic, this
 86 constitutes a model-free, adversarial RL setting requiring policy learning under uncertainty.

87 4 Preliminary solutions

88 We implemented multiple reinforcement learning agents, progressing from tabular methods to
 89 hierarchical extensions, neural approximators, and an LLM-augmented approach. Each model reflects
 90 a different hypothesis about how strategic knowledge in Pokémon battles should be represented and
 91 propagated.

92 4.1 Tabular Q-learning

93 Our baseline tabular Q-learning agent maintains a value table $Q(s, a)$ updated via

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[r_t + \gamma \max_{a'} Q(s_{t+1}, a') - Q(s_t, a_t) \right].$$

94 All variants used linear ϵ -decay from 1.0 to 0.05 and ϵ -greedy action selection. The initial implemen-
 95 tation, restricted to a win/loss reward and no switching, served primarily as an environment sanity
 96 check and achieved negligible performance.

97 Introducing intermediate rewards led to modest improvement (=36% vs. random), but also revealed
 98 reward hacking behaviors (e.g., repeated healing/status moves). Our final version incorporated
 99 switching and additional penalties to suppress these exploits. Despite enlarging the state space
 100 significantly, this model achieved our best results: approximately 98% win rate vs. random (Figure 6)
 101 and 55% vs. maxbp (Figure 5). Additional training and refined shaping could plausibly yield further
 102 gains.

103 4.2 Hierarchical Q-learning

104 To mitigate the combinatorial action space, we decomposed decisions into a master policy (selecting
 105 among four moves and “switch”) and a subpolicy that chooses the replacement Pokémon. This
 106 division drastically reduces the effective Q-table size and stabilizes early learning. As shown in

107 Figure 1, the hierarchical agent quickly surpassed 80% vs. `random`, but required millions of battles
108 to improve further due to persistent exploration demands and slow convergence. Earlier variants
109 with larger state encodings (Figures 2–3) failed to converge even after 2M battles, demonstrating
110 sensitivity to representation choices.

111 4.3 SARSA and Deep Q-network (DQN)

112 Our Linear SARSA and DQN models approximate $Q(s, a)$ in feature space rather than enumerating
113 discrete states. Both used replay buffers, while DQN additionally employed a target network for
114 stabilization. Linear SARSA plateaued at ≈20% against `random` (Figure 9) due to its inability to
115 learn switching behaviors. DQN, however, captured nonlinear patterns and switching strategies,
116 reaching ≈85% win rate vs. `random` (Figure 8). Nonetheless, DQN was highly sensitive to exploration
117 schedules and prone to policy collapse when replay distributions became skewed.

118 4.4 LLM-based reasoning

119 We attempted to integrate the PokéChamp LLM-based opponent-modeling framework, which con-
120 ditions action choices on predicted movepools. However, the system was incompatible with Gen-
121 eration 1 mechanics, and dynamic behaviors such as Ditto transformations repeatedly crashed the
122 opponent-model module. Moreover, running the LLM introduced prohibitive latency (=4 minutes
123 per game) and incurred significant token costs, especially in battles that terminated via the server’s
124 1000-action limit. These constraints prevented systematic evaluation.

125 5 Evaluation

126 We evaluated all agents in the Pokémon Showdown Gen 1 environment using the `poke-env` API.
127 Each model was trained and tested across multiple random seeds, and results were logged to CSV
128 files containing win rates, exploration values, and auxiliary diagnostics. Performance was measured
129 against two standard opponents: the `random` baseline and the rule-based `maxbp` agent.

130 5.1 Tabular Q-Learning

131 Tabular Q-Learning exhibited the strongest performance among all models. Early variants without
132 switching and with sparse rewards plateaued around 36% win rate. Introducing shaped rewards and
133 enabling switching substantially improved strategic depth and led to consistent policy improvement.
134 The final trained agent reached approximately 98% win rate versus `random` and 55–56% versus
135 `maxbp` (Figures 6, 5). These results are notable given that battles involve randomly sampled teams,
136 making some scenarios unwinnable by construction.

137 Despite these gains, the agent remained sensitive to reward design: penalties applied upon switching
138 discouraged long-term advantageous swaps, occasionally biasing the policy toward staying in with
139 unfavorable matchups. A refined reward structure that distinguishes tactical damage from positional
140 disadvantage would likely yield additional improvements.

141 5.2 Hierarchical agent behaviors

142 The hierarchical Q-learning agent quickly achieved strong performance against the `random` opponent,
143 initially exceeding 80% win rate. However, its learning dynamics were unstable: sparse rewards and
144 insufficient exploration caused Q-values to propagate poorly, producing intermittent policy collapse
145 during later training (Figure 1). Variants with larger state encodings failed to converge even after
146 millions of battles, suggesting that while decomposing the action space reduces representational
147 burden, careful reward shaping and slower decay schedules are necessary for stable convergence.

148 5.3 Linear SARSA and DQN results

149 Linear SARSA improved sample efficiency relative to tabular methods, but consistent reward hacking
150 behaviors led to premature convergence toward degenerate policies, often involving repeated moves
151 that maximized intermediate rewards without increasing win rate. Most variants eventually collapsed.

152 DQN was more expressive and successfully learned switching behaviors, reaching win rates of
 153 approximately 85% against random (Figure 8). However, performance proved highly sensitive to
 154 feature design and replay buffer composition. When the distribution of stored transitions became
 155 homogeneous, the policy collapsed, illustrating that deep function approximation in this domain
 156 requires significantly more tuning than tabular methods. Gen 1 Pokéémon is deceptively simple but
 157 structurally hostile to deep RL without dense reward shaping.

Table 1: Performance of implemented agents against baseline opponents in Pokéémon Showdown (Gen 1).

Agent	vs Random	vs MaxBP	Behavior
Tabular Q-Learning	98%	55–56%	Stable, improving
Hierarchical Q-Learning	80–90%	<20%	Partial convergence
Deep Q-Network (DQN)	~85%	~25%	Sensitive, collapses
Linear SARSA	~20%	<5%	No convergence

158 6 Conclusion

159 Our results show that, within the constrained mechanics of Generation 1 Pokéémon, tabular reinforcement
 160 learning remains remarkably competitive. Tabular Q-Learning achieved the best performance,
 161 reaching a 98% win rate against the random agent and 56% against maxbp, outperforming hierar-
 162 chical, linear, and deep variants. While hierarchical methods reduced action-space complexity, they
 163 exhibited unstable convergence, and both DQN and Linear SARSA suffered from reward hacking and
 164 policy collapse. Future work should explore more expressive reward shaping, improved exploration
 165 strategies, and scalable state abstractions to support larger team spaces and later Pokéémon generations.

166 Workload division

- 167 • **Arzaan Singh:** Implemented the Hierarchical Q-Learning, Linear SARSA, and DQN.
- 168 • **Ian Kreger:** Ian setup and handled environment integration with `poke-env`, and developed
 169 the Tabular Q-Learning.
- 170 • **Michael Weild:** Studied the LLM-assisted action reasoning and helped develop the DQN
 171 implementation, running and debugging the model.
- 172 • **All authors:** Participated in experimentation, result interpretation, and the report.

173 References

- 174 Luca Castronovo. poke-env documentation. https://poke-env.readthedocs.io/en/stable/getting_started.html#connecting-bots-to-showdown, 2024. Accessed: 2025-12-06.
- 176 Matteo Dell’Acqua. Alphapoke project delivery. <https://matteoh2o1999.github.io/en-us/projects/alphaPoke-project-delivery>, 2025. Accessed: 2025-12-06.
- 178 Jake Grigsby. Metamon: Reinforcement learning framework for pokéémon. <https://metamon.tech/>, 2024. Accessed: 2025-12-06.
- 180 Akshay Kalose, Kris Kaya, and Alvin Kim. Optimal battle strategy in pokéémon using reinforcement
 181 learning. Technical report, Stanford University, 2018. AA228 final project report, accessed
 182 2025-12-06.
- 183 Sarthak Kannan, Zhibo Chen, Hanxiao Hu, Yiming Chen, Zihan Dai, Haotian Sun, Xiaohui Zhai, and
 184 Yuanzhi Wang. PokéChamp: an expert-level minimax language agent for competitive pokéémon,
 185 2025. URL arxiv.org.
- 186 Seth Karten, Jake Grigsby, Stephanie Milani, Kiran Vodrahalli, Amy Zhang, Fei Fang, Yuke Zhu,
 187 and Chi Jin. The pokeagent challenge: Competitive and long-context learning at scale. In *NeurIPS
 188 Competition Track*, April 2025.

- 189 Guangcong Zarel Luo. Pok  mon showdown. <https://pokemonshowdown.com/>. Accessed: 2025-
190 12-06.
- 191 Brandon Tan and Dhruv Bhatt. Reinforcement learning for pok  mon red. Technical report, Stanford
192 University, 2018. AA228 Final Project Report, Accessed: 2025-12-06.
- 193 Author(s) Unknown. Scalable offline reinforcement learning for pok  mon battles. *arXiv preprint*
194 *arXiv:2503.04094*, 2025. URL <https://arxiv.org/abs/2503.04094>. Accessed: 2025-12-06.
- 195 Jett Meng Wang. Winning at pokemon battles: Reinforcement learning with monte carlo tree search
196 and proximal policy optimization. Master's thesis, Massachusetts Institute of Technology, 2024.
197 Accessed: 2025-12-06.

198 **A Appendix**




Figure 1: Hierarchical Q-learning training metrics: rolling and overall win rates (top), Q-table state growth (middle), and epsilon decay with training speed (bottom).




Figure 2: Hierarchical Q-learning training metrics for First Failed Version Against Random: rolling and overall win rates (top), Q-table state growth (middle), and epsilon decay with training speed (bottom).




Figure 3: Hierarchical Q-learning training metrics for First Failed Version Against MaxBP: rolling and overall win rates (top), Q-table state growth (middle), and epsilon decay with training speed (bottom).




Figure 4: Tabular Q-learning training metrics: Implementation 2 Rolling Win Rate vs Random Player




Figure 5: Tabular Q-learning training metrics: Implementation 3 Overall Win Rate vs Maximum-BasePower Player




Figure 6: Tabular Q-learning training metrics: Implementation 3 Overall Win Rate vs Random Player




Figure 7: DQN training metrics against MaxBasePower: rolling win rate (top left), overall win rate progression (top right), combined win rate comparison (middle), epsilon decay (right middle), and training speed in battles per second (bottom).




Figure 8: DQN training metrics against Random agent: rolling win rate (top left), overall win progression (top right), combined win analysis (middle), epsilon decay (right middle), and training throughput (bottom).




Figure 9: SARSA agent performance against the MaxBP. Left: cumulative win rate showing slow initial convergence followed by consistent improvement once exploration decreases. Right: exponential epsilon decay schedule demonstrating diminishing exploration over training episodes.